

TDS Protocol Documentation

This document attempts to cover the TDS protocol for:

TDS Version	Supported Products
4.2	Sybase SQL Server < 10 and Microsoft SQL Server 6.5
5.0	Sybase SQL Server >= 10
7.0	Microsoft SQL Server 7.0
7.1	Microsoft SQL Server 2000
7.2	Microsoft SQL Server 2005

Contents

- [Common Terms](#)
- [Typical Usage Sequences](#)
- [The Packet Format](#)
- [Login Packet](#)
- [TDS 7.0 Login Packet](#)
- [Collation structure](#)
- [Client requests](#)
- [Server Responses](#)
- [OCBC stored procedures \(by jtids\)](#)

Common Terms

TDS protocol versions

TDS 5.0	tds version 5.0
TDS 7.0	tds version 7.0
TDS 7.0+	tds version 7.0, 7.1 and 7.2
TDS 5.0-	tds version 5.0 and previous

Variable types used in this document:

CHAR	8-bit char
CHAR[6]	string of 6 chars
CHAR[n]	variable length string
XCHAR	single byte (TDS 5.0-) or ucs2le (TDS 7.0+) characters
INT8	8-bit int
INT16	16-bit int
INT32	32-bit int
UCS2LE	Unicode in UCS2LE format

Note: FreeTDS uses TDS_TINYINT for INT8 and TDS_SMALLINT for INT16.

Typical Usage sequences

These are TDS 4.2 and not meant to be 100% correct, but I thought they might be helpful to get an overall view of what goes on.

```
--> Login  
<-- Login acknowledgement
```

```
--> INSERT SQL statement  
<-- Result Set Done
```

```
--> SELECT SQL statement  
<-- Column Names  
<-- Column Info  
<-- Row Result  
<-- Row Result  
<-- Result Set Done
```

```

--> call stored procedure
<-- Column Names
<-- Column Info
<-- Row Result
<-- Row Result
<-- Done Inside Process
<-- Column Names
<-- Column Info
<-- Row Result
<-- Row Result
<-- Done Inside Process
<-- Return Status
<-- Process Done

```

The packet format

Every informations in TDS protocol (query, RPCs, responses and so on) is splitted in packets.

All packets start with the following 8 byte header.

INT8	INT8	INT16	4 bytes
+-----+	+-----+	+-----+	+-----+
packet	last packet	packet	unknown
type	indicator	size	
+-----+	+-----+	+-----+	+-----+

Fields:

packet type

- 0x01 [TDS 4.2 or 7.0 query](#)
- 0x02 [TDS 4.2 or 5.0 login packet](#)
- 0x03 [RPC](#)
- 0x04 [responses from server](#)
- 0x06 cancels
- 0x07 Used in [Bulk Copy](#)

```
0x0F TDS 5.0 query
0x10 TDS 7.0 login packet
0x11 TDS 7.0 authentication packet
0x12 TDS 8 prelogin packet
last packet indicator
0x00 if more packets
0x01 if last packet
packet size
(in network byte order)
unknown?
always 0x00
this has something to do with server to server communication/rpc stuff
```

The remainder of the packet depends on the type of information it is providing. As noted above, packets break down into the types query, login, response, and cancels. Response packets are further split into multiple sub-types denoted by the first byte (a.k.a. the token) following the above header.

Note: A TDS packet that is longer than 512 bytes is split on the 512 byte boundary and the "more packets" bit is set. The full TDS packet is reassembled from its component 512 byte packets with the 8-byte headers stripped out. 512 is the block_size in the login packet, so it could be set to a different values. In Sybase you can configure a range of valid block sizes. TDS 7.0+ use a default of 4096 as block size.

TDS 4.2 & 5.0 Login Packet

Packet type (first byte) is 2. The numbers on the left are decimal offsets *including* the 8 byte packet header.

byte	var	type	description

8	CHAR[30]		host_name
38	INT8		host_name_length
39	CHAR[30]		user_name
69	INT8		user_name_length
70	CHAR[30]		password

```

100  INT8      password_length
101  CHAR[30]  host_process
131  INT8      host_process_length
132  ?        magic1[6]          /* mystery stuff */
138  INT8      bulk_copy
139  ?        magic2[9]          /* mystery stuff */
148  CHAR[30]  app_name
178  INT8      app_name_length
179  CHAR[30]  server_name
209  INT8      server_name_length
210  ?        magic3[1]          /* 0, don't know this one either */
211  INT8      password2_length
212  CHAR[30]  password2
242  CHAR[223] magic4
465  INT8      password2_length_plus2
466  INT16     major_version     /* TDS version */
468  INT16     minor_version     /* TDS version */
470  CHAR      library_name[10]  /* "Ct-Library" or "DB-Library" */
480  INT8      library_length
481  INT16     major_version2    /* program version */
483  INT16     minor_version2    /* program version */
485  ?        magic6[3]          /* ? last two octets are 13 and 17 */
                                   /* bdw reports last two as 12 and 16 here */
                                   /* possibly a bitset flag */
488  CHAR[30]  language          /* e.g. "us-english" */
518  INT8      language_length
519  ?        magic7[1]          /* mystery stuff */
520  INT16     old_secure         /* explanation? */
522  INT8      encrypted         /* 1 means encrypted all password fields blank */
523  ?        magic8[1]          /* no clue... zeros */
524  CHAR      sec_spare[9]       /* explanation? */
533  CHAR[30]  char_set          /* e.g. "iso_1" */
563  INT8      char_set_length
564  INT8      magic9[1]         /* 1 */
565  CHAR[6]   block_size        /* in text */
571  INT8      block_size_length
572  ?        magic10[25]        /* lots of stuff here...no clue */

```

Any help with the magic numbers would be most appreciated.

TDS 7.0+ Login Packet

byte	var	type	description

0	INT32		total packet size
4	INT8[4]		TDS Version
			0x00000070 7.0
			0x01000071 7.1
			0x02000972 7.2 (7.2.9?)
8	INT32		packet size (default 4096)
12	INT8[4]		client program version
16	INT32		PID of client
20	INT32		connection id (usually 0)
24	INT8		option flags 1
			0x80 enable warning messages if SET LANGUAGE issued
			0x40 change to initial database must succeed
			0x20 enable warning messages if USE <database> issued
			0x10 enable BCP
			0x08 use ND5000 floating point format (untested)
			0x04 use VAX floating point format (untested)
			0x02 use EBCDIC encoding (untested)
			0x01 use big-endian byte order (untested)
25	INT8		option flags 2
			0x80 enable domain login security
			0x40 "USER_SERVER - reserved"
			0x20 user type is "DQ login"
			0x10 user type is "replication login"
			0x08 "fCacheConnect"
			0x04 "fTranBoundary"
			0x02 client is an ODBC driver
			0x01 change to initial language must succeed
26	INT8		0x04 spawn user instance (TDS 7.2)
			0x02 XML data type instances are returned as binary XML (TDS 7.2)

```

0x01 password change requested (TDS 7.2)
27  INT8      0x01 SQL Type: 0 = use default, 1 = use T-SQL (TDS 7.2)
28  INT8[4]   time zone (0x88ffffff ???)
32  INT8[4]   collation information
36  INT16     position of client hostname (86)
38  INT16     hostname length
40  INT16     position of username
42  INT16     username length
44  INT16     position of password
46  INT16     password length
48  INT16     position of app name
50  INT16     app name length
52  INT16     position of server name
54  INT16     server name length
56  INT16     position of remote server/password pairs
58  INT16     remote server/password pairs length
60  INT16     position of library name
62  INT16     library name length
64  INT16     position of language
66  INT16     language name (for italian "Italiano", coded UCS2)
68  INT16     position of database name
70  INT16     database name length
72  INT8[6]   MAC address of client
78  INT16     position of auth portion
80  INT16     NT authentication length
82  INT16     next position (same as total packet size)
84  INT16     0
86  UCS2LE[n] hostname
      UCS2LE[n] username
      UCS2LE[n] encrypted password
      UCS2LE[n] app name
      UCS2LE[n] server name
      UCS2LE[n] library name
      UCS2LE[n] language name
      UCS2LE[n] database name
      NT Authentication packet

```

NT Authentication packet

```
0   CHAR[8]    authentication id "NTLMSSP\0"
8   INT32      1 message type
12  INT32      0xb201 flags
16  INT16      domain length
18  INT16      domain length
20  INT32      domain offset
24  INT16      hostname length
26  INT16      hostname length
28  INT32      hostname offset
32  CHAR[n]    hostname
    CHAR[n]    domain
```

See documentation on Samba for detail (or search ntlm authentication for IIS)

For mssql 2005 before hostname (byte 86) you have

```
86  INT16      next position, or
    position of file name for a database to be
    attached during the connection process
88  INT16      database filename length
90  INT16      new password position
92  INT16      new password length
94  UCS2LE[n]  hostname
    ... (as above)
```

"current pos" is the starting byte address for a Unicode string within the packet. The length of that Unicode string immediately follows. That implies there are at least 2 more strings that could be defined. (character set??)

Username and password are empty if domain authentication is used.

If the client uses an authentication packet, the server replies with an [Authentication](#) token followed by an [Authentication packet](#).

TDS 7.0 Authentication Packet


```
varies
+-----+
| auth |
+-----+
```

auth authentication data
for NTLM this message 3

This packet usually follows [Authentication](#) token.

Types

HEX	DEC	type	protocol	nullable	size	collate
0x1F	31	SYBVOID	7+	no	0	
0x22	34	SYBIMAGE		yes	4	
0x23	35	SYBTEXT		yes	4	yes
0x24	36	SYBUNIQUE	7+	yes	1	
0x25	37	SYBVARBINARY		yes	1	
0x26	38	SYBINTN		yes	1	
0x27	39	SYBVARCHAR		yes	1	
0x2D	45	SYBBINARY		yes	1	
0x2F	47	SYBCHAR		yes	1	
0x30	48	SYBINT1		no	0	
0x32	50	SYBBIT		no	0	
0x34	52	SYBINT2		no	0	
0x38	56	SYBINT4		no	0	

0x3A	58	SYBDATETIME4		no	0	
0x3B	59	SYBREAL		no	0	
0x3C	60	SYBMONEY		no	0	
0x3D	61	SYBDATETIME		no	0	
0x3E	62	SYBFLT8		no	0	
0x40	64	SYBSINT1	5	no	0	
0x41	65	SYBUINT2	5	no	0	
0x42	66	SYBUINT4	5	no	0	
0x43	67	SYBUINT8	5	no	0	
0x62	98	SYB VARIANT	7+	yes	4	
0x63	99	SYBNTEXT	7+	yes	4	yes
0x67	103	SYBNVARCHAR	7+	yes	1	
0x68	104	SYBBITN		yes	1	
0x6A	106	SYBDECIMAL		yes	1	
0x6C	108	SYB NUMERIC		yes	1	
0x6D	109	SYBFLT N		yes	1	
0x6E	110	SYBMONEY N		yes	1	
0x6F	111	SYBDATETIM N		yes	1	
0x7A	122	SYBMONEY4		no	0	
0x7F	127	SYBINT8		no	0	
0xA5	165	XSYBVARBINARY	7+	yes	2 *	
0xA7	167	XSYBVARCHAR	7+	yes	2 *	yes
0xAD	173	XSYBBINARY	7+	yes	2	
0xAF	175	XSYBCHAR	7+	yes	2	yes

0xE1	225	SYBLongBinary	5	yes	4	
0xE7	231	XSyBnVarchar	7+	yes	2 *	yes
0xEF	239	XSyBnChar	7+	yes	2	yes

* Under TDS 7.2+ these types allow size to be -1, representing varchar(max), varbinary(max) and nvarchar(max). Data representation for them changes:

- size is 64 (not 16) bits
- size of -1 means NULL
- size of -2 means the size is unknown
- the data are split in chunks, where each chunk starts with a 32-bit size
- a chunk with size <= 0 is the terminal chunk

Collation type - TDS 7.1

The collation structure contains information about the character set encoding and comparison method.

INT16	INT16	INT8
+-----+	+-----+	+-----+
codepage	flags	charset_id
+-----+	+-----+	+-----+

codepage windows codepage (see <http://www.microsoft.com/globaldev/nlsweb/>)
also specified in lcid column of master..syslanguages

flags sort flags

- 0x100 binary compare
- 0x080 width insensitive
- 0x040 Katatype insensitive
- 0x020 accent insensitive
- 0x010 case insensitive

If binary flag is specified other flags are not present

Low nibble of flags is a charset specifier (like chinese dialect)

charset_id charset id in master..syscharsets table or zero for no SQL collations

Collations names can be obtained from select name from ::fn_helpcollations() query

Column Metadata

INT8	XCHAR[n]	INT8	INT32	INT8		
+-----+	+-----+	+-----+	+-----+	+-----+		
column name	column name	flags	user	column		
length			type	type		
+-----+	+-----+	+-----+	+-----+	+-----+		
varies	INT8	INT8	INT16	XCHAR[n]	INT8	varies
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+	+-----+
column size	precision	scale	t length	table name	locale	locale
					length	info
(optional)	(optional)	(optional)	(optional)	(optional)	(opt)	(opt)
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+	+-----+

column name length	
column name	column name in result set, not necessarily db column name
flags	bit flags 0x1 hidden (TDS 5.0) 0x2 key 0x10 writable 0x20 can be NULL 0x40 identity
user type	usertype column from syscolumns
column type	column type
column size	not present for fixed size columns
precision	present only for SYBDECIMAL and SYBNUMERIC
scale	present only for SYBDECIMAL and SYBNUMERIC
t length	present only for SYBTEXT and SYBIMAGE, length of table name
table name	present only for SYBTEXT and SYBIMAGE
locale length	length of locale info (in bytes)

locale info	only for TDS 5.0 results (not for parameters)
	unknown
	only for TDS 5.0 results (not for parameters)

Client request

Normal tokens (contained in packets 0xF)

TODO

Special packets

- 0x1 1 [Language](#)
 - 0x3 3 [RPC](#) TDS 4.6+
 - 0x7 7 [BCP](#) TDS 5.0+
-

Language packet (0x1 1)


This sample packet contain just SQL commands. It's supported by all TDS version (although TDS 5.0 have others token with similar use)

```
XCHAR[n]
+-----+
| string |
+-----+
```

string SQL text

RPC packet (0x3 3)

Do not confuse an RPC packet with an RPC token. The RPC packet is supported by all version of TDS; the RPC token is supported only by TDS 5.0 (and has different format). This is the oldest (and the only one in mssql) way to call directly an RPC. Sybase also documents it, but as 0xE.



name length length of RPC name in characters.
mssql2k+ support some core RPC using numbers
If a number is used instead of name name length is marked as -1 (null) and a INT16 is used for the name.

0x1	1	sp_cursor
0x2	2	sp_cursoropen
0x3	3	sp_cursorprepare
0x4	4	sp_cursorexecute
0x5	5	sp_cursorprepexec
0x6	6	sp_cursorunprepare
0x7	7	sp_cursorfetch
0x8	8	sp_cursoroption
0x9	9	sp_cursorclose
0xA	10	sp_executesql
0xB	11	sp_prepare
0xC	12	sp_execute ???
0xD	13	sp_prepexec
0xE	14	sp_prepexecrpc
0xF	15	sp_unprepare

sp_execute seems to have some problems, even MS ODBC use name version instead of number.

rpc name name of RPC.

flags bit flags.

0x1	1	recompile procedure (TDS 7.0+/TDS 5.0)
0x2	2	no metadata (TDS 7.0+)

(I don't know meaning of "no metadata" -- freddy77)

params parameters. See below

Every parameter has the following structure

```
+-----+-----+
| data info | data |
+-----+-----+
data info    data information. See below
data        data. See results for detail
```

Data info structure

```
      INT8      XCHAR[n]      INT8      INT32
+-----+-----+-----+-----+
| name length | param name | flags | usertype (TDS 5.0) |
+-----+-----+-----+-----+
      INT8  varies  varies      INT8[5]      INT8
+-----+-----+-----+-----+-----+
| type | size | optional | collate      | locale      |
|      | (opt) | (opt)   | info(TDS 7.1) | length (TDS 5.0) |
+-----+-----+-----+-----+-----+
```

name length parameter name length (0 if unused)

param name parameter name

flags bit Name Meaning

0x1 TDS_RPC_OUTPUT output parameter

0x2 TDS_RPC_NODEF output parameter has no default value.
Valid only with TDS_RPC_OUTPUT.

usertype usertype

type param type


size see [Results](#)

optional see [Results](#). Blobs DO NOT have
optional on input parameters (output blob parameters
are not supported by any version of TDS).

collate info only for type that want collate info and using TDS 7.1
locale length locale information length. Usually 0 (if not locale
information follow, the structure is unknown)

Chained RPCs


Under TDS 7.0+ is possible to chain multiple RPCs together. This is useful to limit packets and round-trips with server. RPCs can be chained using byte 0x80 (TDS 7.0/TDS 7.1) or 0xFF (TDS 7.2).



Bulk Copy packet (0x7 7)

This documents a TDS 5.0 packet. It might be true for others....

BCP Packet Structure



followed by N row buffers,
where N is computed by exhausting the packet size

BCP Packet Row Buffer

INT16	INT8	INT8	INT16
size	ncols	zero	size (again)

followed by column buffers (data), where

first, the fixed-size datatype columns

fixed size and count (determined by column definition)

data	[repeats once for each mandatory column]
-----------	--

then, the variable-size (including nullable) datatype columns

variable size and count

data	[repeats ncol times]
-----------	----------------------

followed by two tables (!) to describe the column buffers

Adjustment Table (optional)

INT8	INT8
1 + ncols	offset

[repeats ncol + 1 times]

Offset Table (mandatory)

INT8	INT8
1 + ncols	offset

[repeats ncol + 1 times]

The BCP packet has a slightly different Packet header!?

Computation of Offset and Adjustment tables

The offset and adjustment tables describe the position of the first byte of each variable-size column. The first element holds the count of elements in the offset/adjustment table. Thereafter, the offsets are arranged in reverse order: the last element — which is also the last byte of the row buffer — holds the offset from the start of the row of the first variable-size column. The next-to-last offset table element holds the starting position of the second variable-size column, and so on.

Offset Table Example

1. 5
2. 31
3. 22
4. 21
5. 8
6. 4

The first element is 5 because there are five elements in the list. There are 4 column data buffers with 5 endpoints. The first column's data begins at offset 4. Computations:

column 1

offset 4

length: $4 = 8 - 4$

column 2

offset 8

length: $13 = 21 - 8$

column 3

offset 21

length: 1 = 22 - 21
column 4
offset 22
length: 9 = 31 - 22

Any column not accounted for is implicitly NULL. To represent a NULL column between two dataful columns, the offset table will have adjacent entries of the same value.

Adjustment table

The so-called adjustment table provides for longer rows. The reader will note the Offset table has 8-bit elements, which would limit the width of the table: the last variable column would have to end less than 256 bytes from the start of the row. Rather than changing the definition of the Offset table, a second table, the Adjustment table, was introduced. It holds high-order bytes for the column offsets.

In other words, to compute a variable column's offset from the start of the row buffer, the server looks up its offset table value, then consults the same position in the adjustment table, and splices them together.

Offset table commentary

The BCP packet is very dense. The data formats are governed by the table definition. Non-NULL columns of course must be present; there is no need to count them or compute their size. The NULL columns are unlimited; their boundaries are defined by the minimalist offset table.

The Adjustment table seems silly at first glance. Why not just make the offset table's elements 16 or even 32 bits? The reason is overhead. Most rows will have less than 256 bytes of variable column data. By using the adjustment table, the BCP packet avoids adding one or even three empty bytes per column per row.

Why is the table in reverse order? Because that places the first offset at a known location: the end of each *row* gives the start of the first column. The server can work its way down the offset table and compute the column sizes. If they

don't add up — if there are data between the (presumed) end of the last column and the start of the offset table — the server knows it should look for an adjustment table. Because the scheme is infinitely repeatable, rows could one day grow to terabyte widths without redefining the packet structure.

Server Responses

Responses from the server start with a single octet (token) identifying its type. If variable length, they generally have the length as the second and third bytes

Tokens encountered thus far:

HEX	DEC	name	note
0x20	32	Param Format 2	5.0 only
0x21	33	Language	5.0 only, client-side
0x22	34	OrderBy 2	5.0 only??
0x61	97	Row Format 2	5.0 only
0x71	113	"Logout"	5.0? ct_close(), client-side?
0x79	121	Return Status	
0x7C	124	Process ID	4.2 only
0x80	128	Cursor Close	5.0 only
0x81	129	Cursor Delete	5.0 only
0x81	129	7.0 Result	7.0 only
0x82	130	Cursor Fetch	5.0 only
0x83	131	Cursor Info	5.0 only
0x84	132	Cursor Open	5.0 only

0x86	134	Cursor Declare	5.0 only
0x88	136	7.0 Compute Result	7.0 only
0xA0	160	Column Name	4.2 only
0xA1	161	Column Format	4.2 only
0xA3	163	Dynamic 2	5.0 only
0xA4	164	Table names	name of tables in a FOR BROWSE select
0xA5	165	Column Info	column information in a FOR BROWSE select
0xA6	166	Option Cmd	5.0 only
0xA7	167	Compute Names	
0xA8	168	Compute Result	
0xA9	169	Order By	
0xAA	170	Error Message	
0xAB	171	Info Message	
0xAC	172	Output Parameters	
0xAD	173	Login Acknowledgement	
0xAE	174	Control	
0xD1	209	Data --- Row Result	
0xD3	211	Data --- Compute Result	
0xD7	215	Params	5.0 only
0xE2	226	Capability	5.0 only. Information on server
0xE3	227	Environment Change	(database change, packet size, etc...)
0xE5	229	Extended Error Message	
0xE6	230	DBRPC	5.0 only RPC calls
0xE7	231	Dynamic	5.0 only

0xEC 236	Param Format	5.0 only
0xED 237	Authentication	7.0 only
0xEE 238	Result Set	5.0 only
0xFD 253	Result Set Done	
0xFE 254	Process Done	
0xFF 255	Done inside Process	

Param Format 2 - TDS 5.0 (0x20 32)

TODO.

Language - TDS 5.0 (0x21 33)

INT32	INT8	CHAR[n]
+-----+	+-----+	+-----+
length	status	query
+-----+	+-----+	+-----+
length	total token length	
status	0 no args	
	1 has args (followed by PARAMFMT/PARAMS)	
query	query (total length - 1)	

Order By 2 (0x22 34)

TODO.

Row Format 2 - TDS 5.0 (0x61 97)

TODO.

"Logout" (0x71 113)

No information. (1 byte, value=0 ?)

Return Status (0x79 121)

INT32

```
+-----+  
| Return status |  
+-----+
```

The return value of a stored procedure.

Process ID (0x7C 124)

8 bytes

```
+-----+  
| process number |  
+-----+
```

Presumably the process ID number for an executing stored procedure. (I'm not sure how this would ever be used by a client. *mjs*)

Cursor Close - TDS 5.0 (0x80 128)

TODO.

Cursor Delete - TDS 5.0 (0x81 129)

TODO.

Result - TDS 7.0+ (0x81 129)

INT16
+-----+-----+
| #columns | column_info |
+-----+-----+

The TDS 7.0 column_info is formatted as follows for each column:

INT16	INT16	INT8	varies	varies	INT8[5]	INT8	UCS2LE[n]
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+	+-----+	+-----+
usertype	flags	type	size	optional	collate	name length	name
			(opt)	(opt)	info(TDS 7.1)		
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+	+-----+	+-----+

usertype

flags

type

size

type modifier

bit flags

0x1 can be NULL

0x8 can be written (it's not an expression)

0x10 identity

data type, values >128 indicate a large type

none for fixed size types

4 bytes for blob and text
2 bytes for large types
1 byte for all others

optional

	INT8	INT8
	+-----+	+-----+
numeric/decimal types:	precision	scale
	+-----+	+-----+
	INT16	UCS2LE[n]
	+-----+	+-----+
blob/text types:	table name length	table name
	+-----+	+-----+

collate info are available only using TDS 7.1 and for characters types (but not for old type like short VARCHAR, only 2byte length versions)

Cursor Fetch - TDS 5.0 (0x82 130)

TODO.

Cursor Info - TDS 5.0 (0x83 131)

TODO.

Cursor Open - TDS 5.0 (0x84 132)

TODO.

Cursor Declare - TDS 5.0 (0x86 134)

TODO.

Compute Result - TDS 7.0+ (0x88 136)

TODO.

Column Name (0xA0 160)

INT16		INT8	CHAR[n]		INT8	CHAR[n]
+-----+		+-----+	+-----+		+-----+	+-----+
total length		length1	column1 name	lengthN	columnN name
+-----+		+-----+	+-----+		+-----+	+-----+

This token is the first token that contain result informations. Is usually followed by [Column Format](#) token (0xA1 161)

Column Format (0xA1 161)

INT16	
+-----+	
total length	column_info
+-----+	

The number of columns is the same of previous [Column Name](#) token.

The TDS 4.2 column_info is formatted as follows for each column:

INT8[4]		INT8	varies	varies
+-----+		+-----+	+-----+	+-----+
usertype/	type	size	optional	
flags		(opt)	(opt)	
+-----+		+-----+	+-----+	+-----+

usertype/flags for Sybase

```
INT32
+-----+
| usertype |
+-----+
```

usertype/flags for MSSQL

```
INT16
+-----+-----+
| usertype | flags |
+-----+-----+
```

usertype type modifier
flags bit flags (only MSSQL)
 0x1 can be NULL
 0x8 can be written (it's not an expression)
 0x10 identity
[type](#) data type
size none for fixed size types
 4 bytes for blob and text
 1 byte for all others
 (TDS 4.2 do not support large types)

optional

```
                    INT8            INT8
numeric/decimal types:    +-----+-----+
(supported??)            | precision | scale |
                    +-----+-----+
```

```
                    INT16                    CHAR[n]
blob/text types:        +-----+-----+
                    | table name length | table name |
                    +-----+-----+
```

Dynamic 2 - TDS 5.0 (0xA3 163)

TODO.

Option Cmd - TDS 5.0 (0xA6 166)

TODO.

Compute Result (0xA8 168)

INT16	INT16	INT8	varies	INT8	INT8[n]
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+
total length	compute id	#columns	column info	#bycols	bycol
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+

column info:

INT8	INT8	INT32	INT8	varies	INT8	varies
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+	+-----+
operator	operand	usertype	column	size	locale length	locale info
			type	(opt)	info (TDS 5.0)	(TDS 5.0)
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+	+-----+

operator	operator
	0x4b COUNT
	0x4c UNSIGNED? COUNT
	0x4d SUM
	0x4e UNSIGNED? SUM
	0x4f AVG
	0x50 UNSIGNED? AVG
	0x51 MIN
	0x52 MAX
	0x09 COUNT_BIG (mssql2k)

	0x30	STDEV (mssql2k)
	0x31	STDEVP (mssql2k)
	0x32	VAR (mssql2k)
	0x33	VARP (mssql2k)
	0x72	CHECKSUM_AGG (mssql2k)
operand	???	
usertype	usertype	
<u>column type</u>	data type	
size	data size	
locale length	length of locale informations	
locale info	locale informations (unknown)	

Each bycol information contains column info for a specific column.

TODO: optional possible?? collate infos ??

TabName (0xA4 164)

TDS4/5/7:

INT16	INT16	XCHAR[n]	
+-----+	+-----+	+-----+	+
total length	name length	table name	...
+-----+	+-----+	+-----+	+
name length	table name length		
table name	table name		

TDS 7.1:

INT16	varies
+-----+	+-----+
total length	table names
+-----+	+-----+

```
table name:
  INT8          INT16          XCHAR[n]
+-----+-----+-----+
| # of components | component length | component name |
+-----+-----+-----+

ie:
  name          -> 01  04 00  ucs2le "name"
  db..name      -> 03  02 00  ucs2le "db"  00 00  04 00  ucs2le "name"
  db.dbo.name   -> 03  02 00  ucs2le "db"  03 00  ucs2le "dbo"  04 00  ucs2le "name"
```

Column Info (0xA5 165)

```
  INT16          varies
+-----+-----+
| total length | column infos |
+-----+-----+
```

column info:

```
  INT8  INT8          INT8  INT8          XCHAR[n]
+-----+-----+-----+-----+-----+
| index | table index | flags | name length | column name |
|       |             |      | (opt)      | (opt)      |
+-----+-----+-----+-----+-----+
```

- index index in result format (1-based)
- table index index in previous TabName (1-based)
- 0 means no table (ie computed)
- flags set of flags
 - 0x04 expression
 - 0x08 key
 - 0x10 hidden
 - 0x20 column name present
- name length length of following column
- column name real column name (result contain the label)

This token follow TabName token

compute "control" ? (0xA7 167)

"control" (0xAE 174)

Miscellaneous note (from *bdw* ?) found with 0xAE:

```
has one byte for each column,  
comes between result(238) and first row(209),  
I believe computed column info is stored here, need to investigate
```

Order By (0xA9 169)

```
INT16      variable (1 byte per col)  
+-----+-----+  
| length | orders |  
+-----+-----+
```

length	Length of packet(and number of cols)
orders	one byte per order by indicating the column # in the output matching the order from Column Info and Column Names and data in following Row Data items. A 0 indicates the column is not in the resulting rows.

an example:

```
select first_name, last_name, number from employee  
order by salary, number  
assuming the columns are returned in the order
```

queried:
first_name then last_name, then number. we would have:


| 2 | 0 | 3 |

where length = 2 then the orders evaluate:
0 for salary, meaning there is no salary data returned
3 for number, meaning the 3rd data item corresponding
to a column is the number

Error Message (0xAA 170)

Non-error Message (0xAB 171)

Extended Error Message (0xE5 229)



length	Length of packet
msg number	SQL message number
state	?
level	An error if level > 10, a message if level <= 10

m length	Length of message
message	Text of error/message
s length	Length of server name
server	Name of "server" ?
p length	Length of process name
process name	Stored procedure name, if any
line#	Line number of input which generated the message

Output Parameters (0xAC 172)

Output parameters of a stored procedure.

INT16	INT8	XCHAR[n]	INT8	INT32	INT8	
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+	+-----+
length	c length	colname	flags	usertype	datatype
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+	+-----+

length	Length of packet
c length	Length of colname
colname	Name of column
flags	0x1 Nullable
usertype	cf. systypes table in database
datatype	Type of data returned

The trailing information depends on whether the datatype is a fixed size datatype.

	N bytes
	+-----+
Datatype of fixed size N	data
	+-----+

	INT8	INT8	N bytes
	+-----+	+-----+	+-----+
Otherwise	column size	actual size N	data
	+-----+	+-----+	+-----+

Login Acknowledgement (0xAD 173)

INT16	INT8	4 bytes	INT8	XCHAR[n]	4 bytes
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+
length	ack	version	t length	text	ser_ver
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+
length	length of packet				
ack	0x01 success 4.2				
	0x05 success 5.0				
	0x06 failure 5.0				
version	TDS version 4 bytes: major.minor.?.?				
t length	length of text				
text	server name (ie 'Microsoft SQL Server')				
ser_ver	Server version				
	(with strange encoding, differring from TDS version)				

Data - Row Result (0xD1 209)

Data - Compute Result (0xD3 211)

INT8	variable size
+-----+	+-----+
token	row data
+-----+	+-----+

Row data starts with one byte (decimal 209), for variable length types, a one byte length field precedes the data, for fixed length records just the data appears.

Note: nullable integers and floats are variable length.

For example: sp_who

The first field is spid, a smallint

The second field is status a char(12), in our example "recv sleep "

The row would look like this:

byte 0 is the token
bytes 1-2 are a smallint in low-endian
byte 3 is the length of the char field
bytes 4-15 is the char field

byte	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15			
hex	D1	01	00	0C	72	65	63	76	20	73	6C	65	65	70	20	20			
	209	1	0	12	r	e	c	v	'	'	s	l	e	e	p	'	'	'	'

Params - TDS 5.0 (0xD7 215)

TODO.

Capability - TDS 5.0 (0xE2 226)

INT16	variable
+-----+-----+	
length	capabilities
+-----+-----+	
length	Length of capability string
capabilities	Server capabilities? Related to login magic?

Environment change (0xE3 227)

INT16	INT8	INT8	CHAR[n]	INT8	CHAR[n]
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+
length	env code	t1 length	text1	t2 length	text2
+-----+	+-----+	+-----+	+-----+	+-----+	+-----+

env code	Code for what part of environment changed
0x01	database context
0x02	language
0x03	character set
0x04	packet size
0x05	TDS 7.0+ LCID
0x06	TDS 7.0+ ??? (sort method? sql server encoding?)
0x07	Collation info
text1	Old value
text2	New value

text1 and text2 are text information (coded in ucs2 in TDS 7.0+) except collation info that's a structure (see collation structure)

DBRPC - TDS 5.0 (0xE6 230)

TODO.

Dynamic - TDS 5.0 (0xE7 231)

TODO.

Param Format - TDS 5.0 (0xEC 236)

INT16	INT16	variable size
-------	-------	---------------

length	number of parameters	parameter info

length length of message following this field
number of parameters number of parameter formats following
list of formats I (*bdw*) imagine it uses the column format structure.

Authentication - TDS 7.0 (0xED 237)

length	auth

length length of authentication data following this field
auth authentication data
for NTLM this is message 2

Client reply with [Authentication packet](#).

Result Set - TDS 5.0 (0xEE 238)

length	number of columns	column info

Fields:
length length of message following this field
number of columns number of columns in the result set, this many column

column info

information fields will follow.
[column info](#)

Done Packets

- Result Set Done (0xFD 253)
- Process Done (0xFE 254)
- Done Inside Process (0xFF 255)

INT16	INT16	INT32	INT64
+-----+	+-----+	+-----+	+-----+
bit flags	unknown	row count (TDS 7.1-)	row count (TDS 7.2)
+-----+	+-----+	+-----+	+-----+

Fields:	
bit flags	0x01 more results 0x02 error (like invalid sql syntax) 0x10 row count is valid 0x20 cancelled
unknown	2,0 /* something to do with block size perhaps */
row count	number of rows affected / returned in the result set. row count is 64-bit using TDS 7.2. (FIXME check if "affected / returned" is correct)

"Result Set Complete" is the end of a query that doesn't create a process on the server. I.e., it doesn't call a stored procedure.

"Process Done" is the end of a stored procedure

"Done In Process" means that a query internal to a stored procedure has finished, but the stored procedure isn't done overall.

Acknowledgements

The following people have contributed to this document:

Brian Bruns (first draft, protocol discovery)

Brian Wheeler (protocol discovery)

Mark Schaal (second draft)

Frediano Ziglio

(short list)

Document Status

\$Id: tds.html,v 1.41 2008-11-25 23:38:33 jklowden Exp \$

